

Global States :-

1000

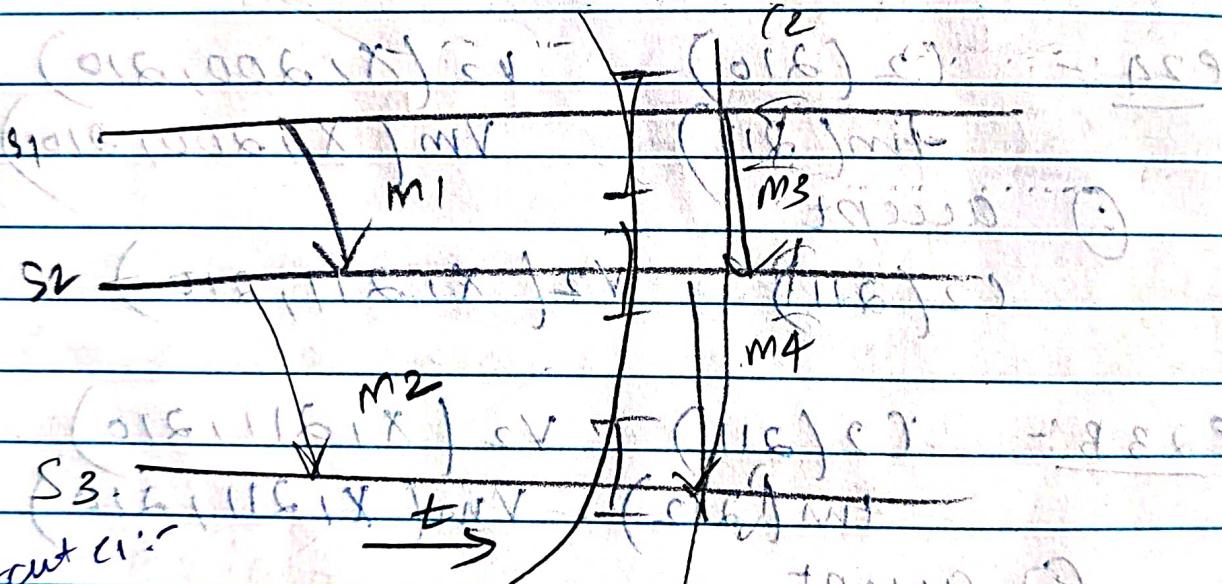
ATM

Bank
Server

while withdrawing, bank is secretly take
states

(union of) all local states = global state
has send & receive events.

$$GS = \bigcup_{i=1}^n LS_i \text{ if and only if } \text{ send } + \text{ receive}$$



$$LS_1 = \{ \text{send}(m_1) \} \cup \{ \dots \}$$

$$LS_2 = \{ \text{recv}(m_2), \text{recv}(m_3) \}.$$

$$LS_3 = \{ \text{recv}(m_4) \}.$$

If every send has its receive \rightarrow
Strongly consistent GS.
consistent gs

Cut & graphical line drawn on timeline.

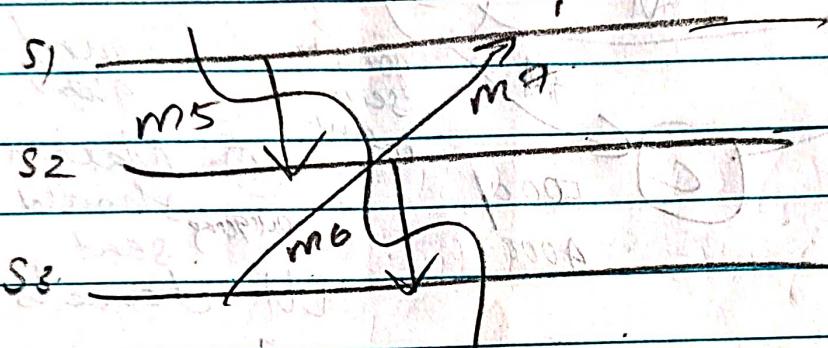
For cut c2:

$$\left. \begin{array}{l} LS_1 = \{ \text{send}(m_3) \} \\ LS_2 = \{ \text{send}(m_4) \} \\ LS_3 = \{ \text{rec}(m_4) \} \end{array} \right\} \text{no strongly consistent}$$

Every Recv should have corresponding send \rightarrow CONSISTENT BS.

It can have transits.

\hookrightarrow only the send() messages is recorded, but not their rec()



$LS_1 = \{ \cdot \}$ \Rightarrow S1 is sending msg to everyone but none is receiving it

$$\left. \begin{array}{l} LS_2 = \{ \text{rec}(m_5) \} \\ LS_3 = \{ \text{rec}(m_6) \} \end{array} \right\} \text{INCONSISTENT}$$

\hookrightarrow no send() for receive message()
[receive has been recorded before send is recorded]

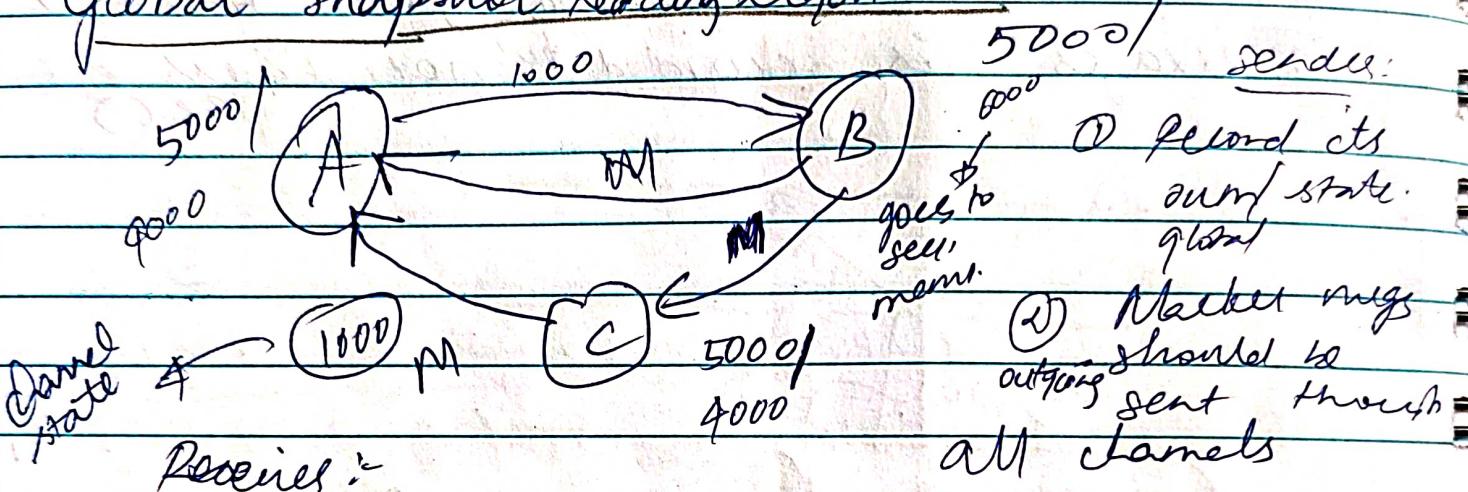
$$LS_3 = \{ \text{send}(m_6), \text{rec}(m_7) \}$$



\hookrightarrow there is a crash, no evidence of send is true.

packet is sent,
but there is a
crash detection.
So the msg is
infinitely in transit.

Global snapshot Receiving Algorithm :-



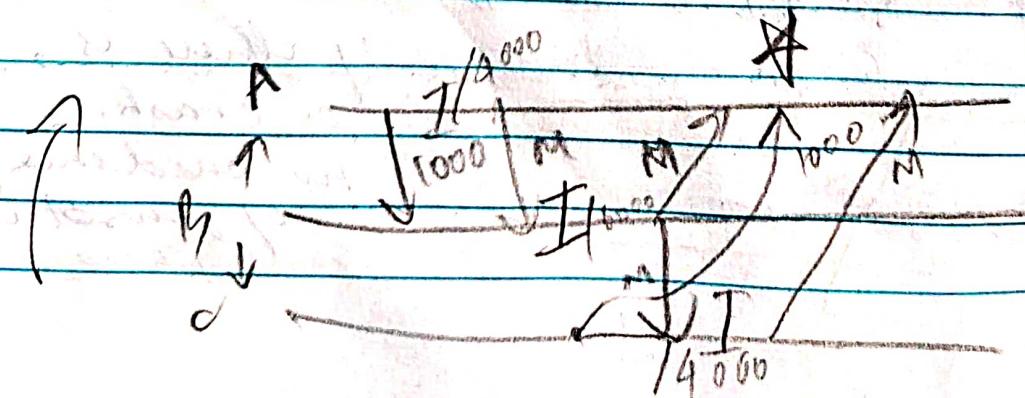
Received :-

Snapshot marker & pass the marker to other outgoing channels (changes as secede)

C stops when it has seen the marker msg from all incoming channels

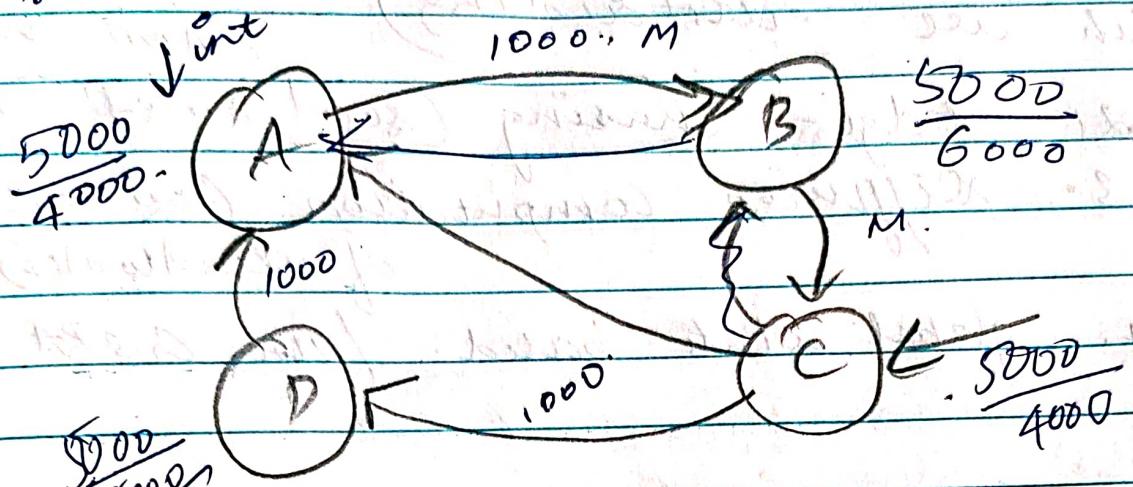
Recieve of B to C it snapshoted but not yet (∴ inconsistent) Thus schedule the msg

Channel state :-



If any message has been wanted in channel incoming, then the msg + snapshot

If snapshot is already taken in the msg in channel state is considered as 1st trans'd.



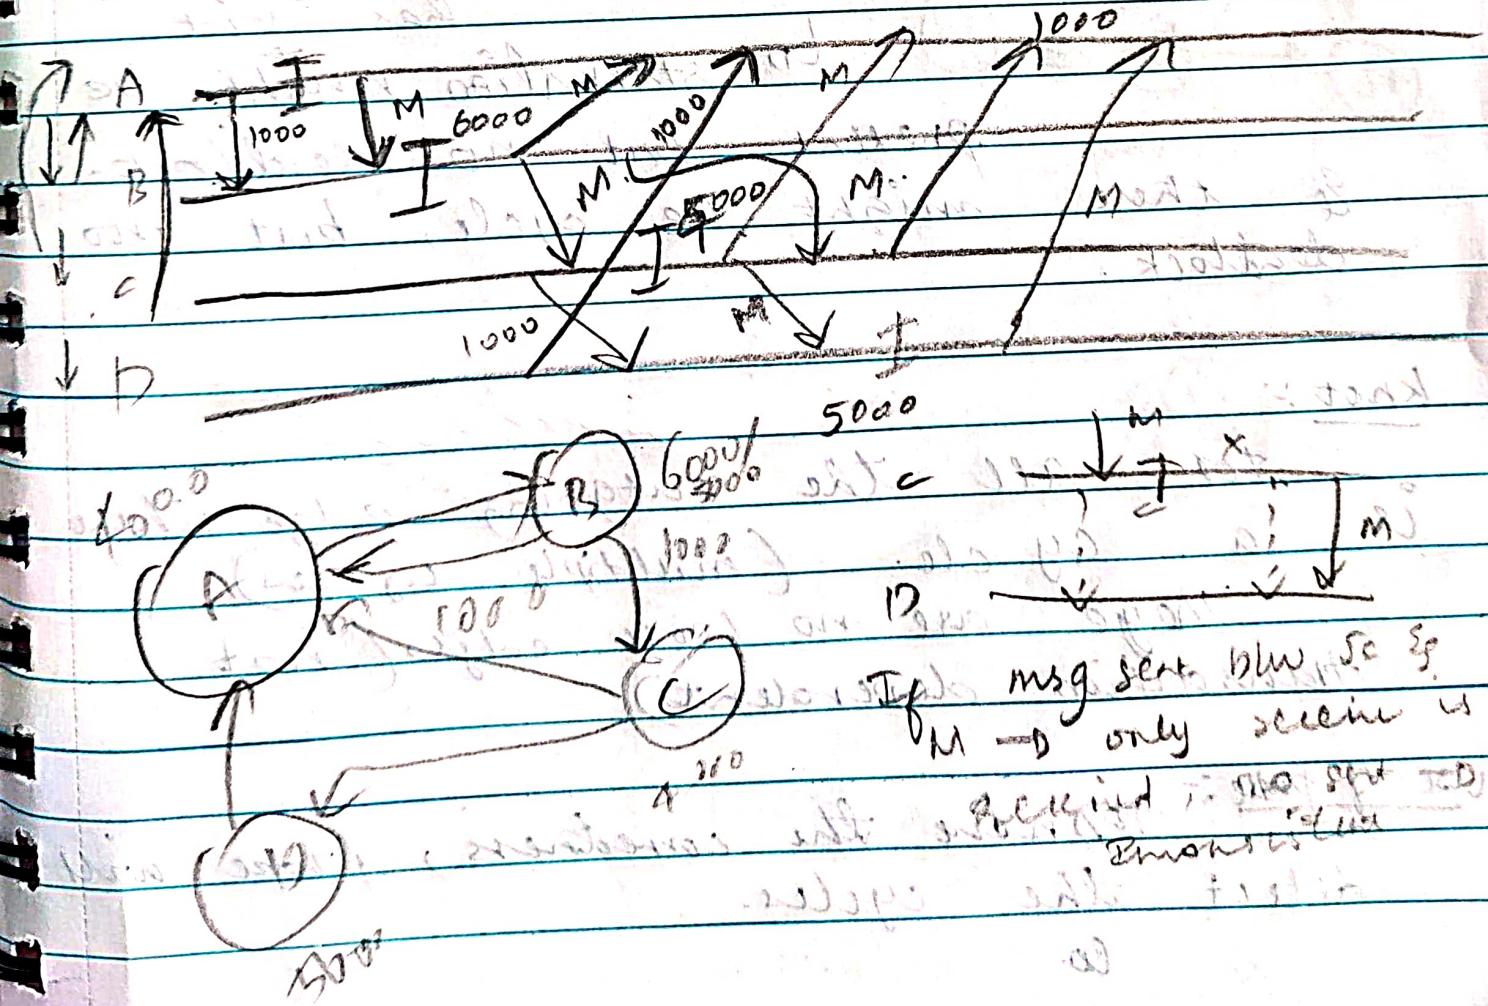
$t_1 := \text{init. } A \rightarrow B.$

$t_2 := \text{init.}$

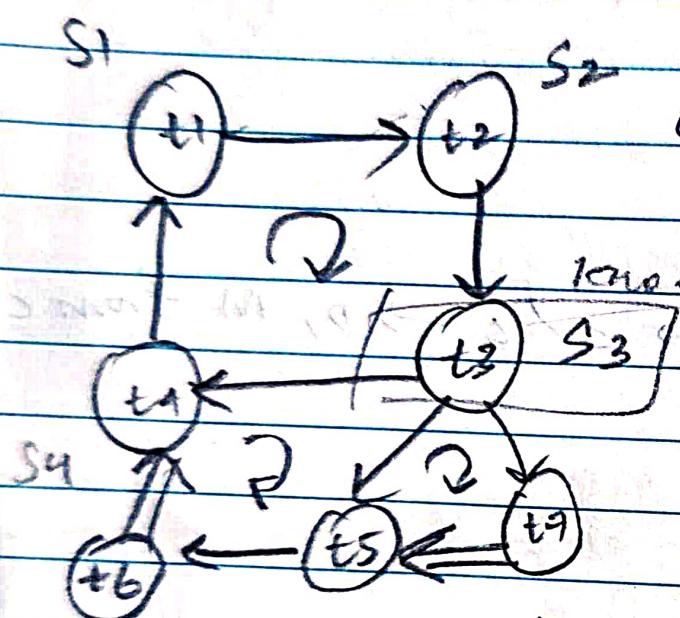
$t_3 := D \rightarrow A$

$t_4 := C \rightarrow A, B \rightarrow C / C \rightarrow D, M \text{ from } c$

$t_5 := C \rightarrow A.$



- classification of deadlock (Knaypp's) :-
1. path pushing (push all the edges which are outstanding) (AND model)
 2. Edge-chasing (send small probe msgs)
 3. Diffusion computation (OR kind of deadlocks)
 4. global state based. (use GORA algo)



OR Request Model

A system can continue if it gets even one resource no matter how many replies it has sent.

No starvation might be present, but no deadlock.

There might be cycle, but no deadlock.

knot:

for all the outgoing edge there is a cycle. (multiple cycles)

There are no loop edges (not transaction dependent)

use of probe:- To prove the correctness, probe will detect the cycles.

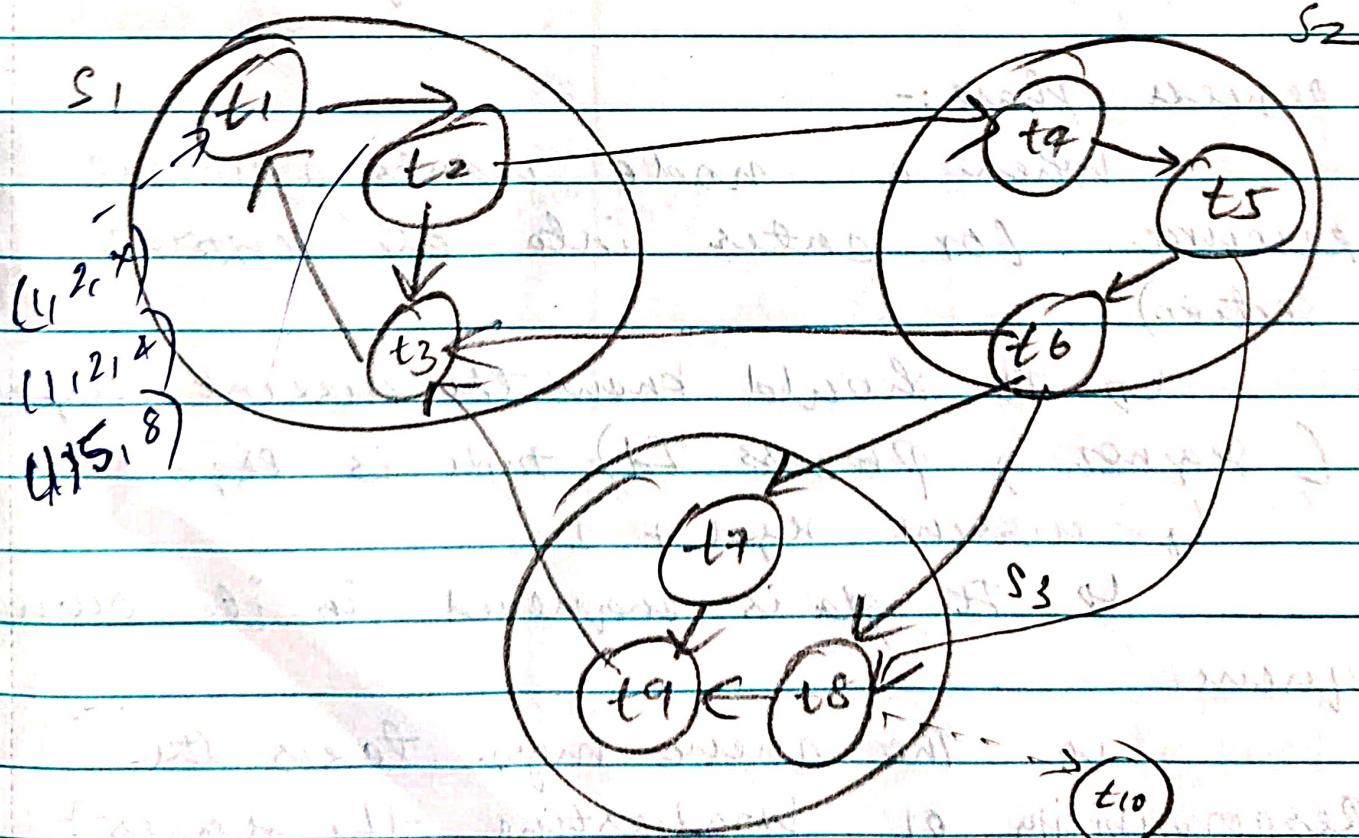
Initiator sends a query if ~~none~~ (i,j,k) is
query is sent only when there is an outgoing
edge.

If the

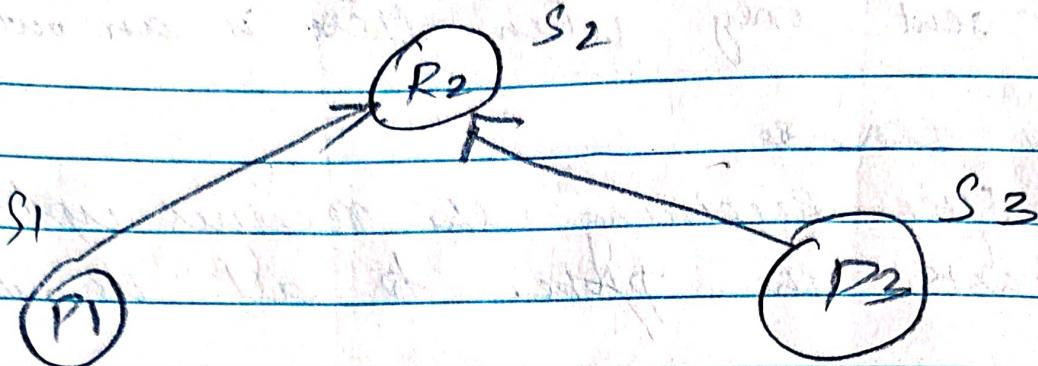
After Receiving the receiver updates &
broadcasts the probe to all its outgoing
edge.

All the nodes, after Receiving has to send
back a reply to the sender after
it receives a reply when it sends
to another node.

- ↳ only when a transaction is waiting,
it receives a reply
- ↳ only when it receives Reply for
all queries, it replies.



DISTRIBUTED MUTUAL EXCLUSION:



- ① Non token Based (Queue Distributed)
- ② token Based:

Lamport

Request

(forcefully
maintain
consistency)

Lamport's Algo:

(non token) HAS 3 phases but it should follow the order in which they are requested

↳ Request, Reply, Release

Request phase :-

When a node wants a critical resource. (or enter into the critical section)

↳ It should know the current request (seg no., process Id) tuple is reqd.

↳ current reqst + 1

↳ Its place is enqueued in its own queue

↳ The queue mgr. takes the responsibility of broadcasting the request to all. (consistency is forcefully maintained by this)

Reply :- Everyone responds to the request in order in which they get Request Condition :

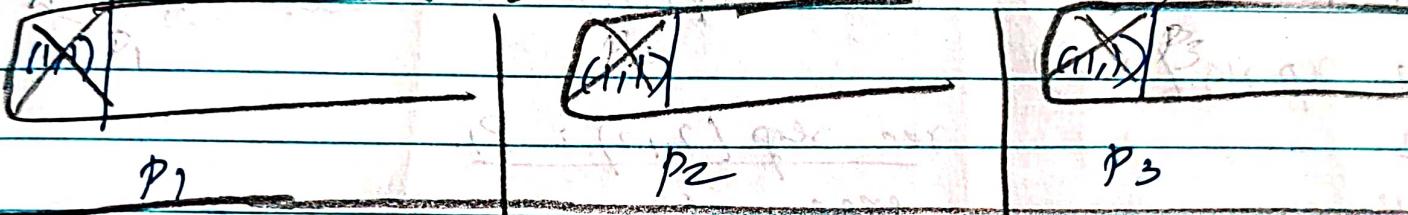
If 2 process are replying to same resource

Release :-

After use, release the resource, so that others can use.

Reliable channel is reqd.

1 both P2 & P3 are not accessing resource



1. Req (1,1)

2.

3.

4. rec reply (1,1) : P2

5. rec reply (1,1) : P3

6. exec CS

7. sel req (1,1) (deletes its req)

8.

rec req (1,1)

del req (1,1)

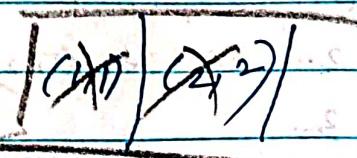
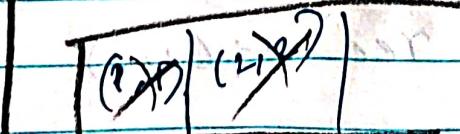
rec req (1,1)

Rep req (1,1)

rec sel (1,1)

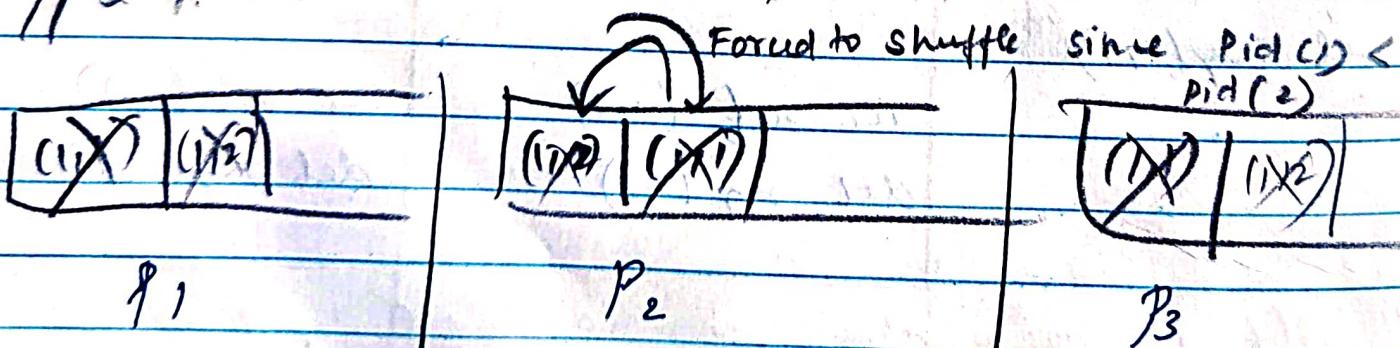
del req (1,1)

11 b P2 is also interested in accessing the Resource



1. seq(1,1)		dec seq(1,1)
2.		seq(2,2)
3.		rep seq(1,1)
4. rec seq(1,1) : P ₃		dec seq(2,2)
5. rec rep(1,1) : P ₂		rep seq(2,2)
6. dec seq(2,2)	<u>dec rep(2,2): P₃</u>	
7. exec CS.		
8. Rel seq(1,1)		
9.	dec del(1,1)	dec del(1,1)
10.	del seq(1,1)	del seq(1,1)
11. rep seq(2,2)		
12.	<u>rec seq(2,2): P₁</u>	
13.	exec CS	
14.	rel seq(2,2)	
15: rec del(2,2)		rec del(2,2)
{ 16. del seq(2,2)		del seq(2,2)}

11.2 process requests for CR at same logical time



- | | |
|-----------------|--------------|
| 1. seq(1,1) | seq(1,2) |
| 2. dec seq(1,2) | rec seq(1,1) |
| 3. | dec seq(1,1) |
| | rec seq(4,2) |

4.		rep reg (1, 1)
5.		rep reg (1, 2)
6.	<u>rec rep(1, 1) : P₃</u>	<u>rec rep(1, 2) : P₃</u>
11.	STARVATION	HAPPENS HERE
7.		rep reg (1, 1)
8.	<u>dec rep(1, 1) : P₂</u>	
9.	exec CS	
10.	del reg(1, 1)	
11.		dec sel(1, 1)
		del reg(1, 1)
12.	rep reg(1, 2)	
13.		rec (rep(1, 2) : P ₁)
14.		exec CS
15.		del reg(1, 2)
16.	rec ref(1, 2)	rec sel(1, 2)
17.	del reg(1, 2)	del reg(1, 2)

Constraint for bringing Fairness :-

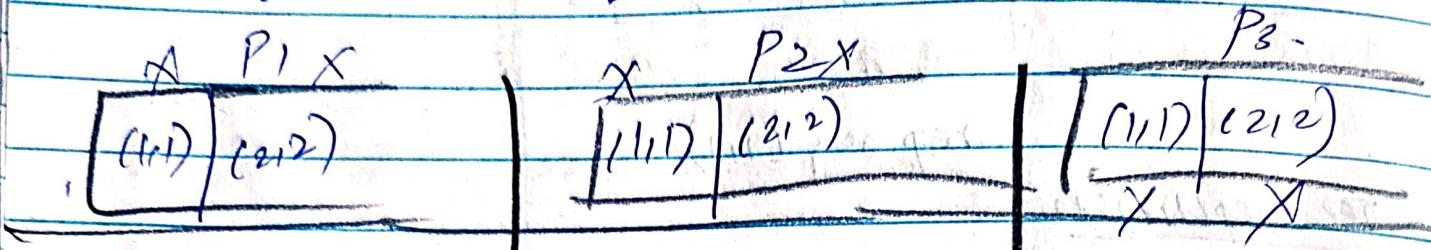
① For a process, if one of the req is outstanding, cannot bring another req for same process in the queue

② If the current req no. is greater than the received req no, then the received request is old, so delete it & continue.

Total no. of msgs. reqd = $\boxed{3 * (n-1)}$
 $(n-1)$ req, $(n-1)$ rep, $(n-1)$ sel for 'n' processes.

Ricart - Agarwala's Algo (D-Mutex) (non-token)

Message complexity is reduced to $2(N-1)$



1. Request (1,1)

Recv Req (2,2)

Recv Rep (1,1) : P2

Recv Rep (1,1) : P3

exec C.S

Rep Req (2,2)

Recv Req (1,1)

Request (2,2)

Rep Rep (1,1)

Recv Rep (2,2) : P3

Revs

Recv V Reply (2,2) : P1

exec C.S

Recv Req (1,1)

Rep Rep (1,1)

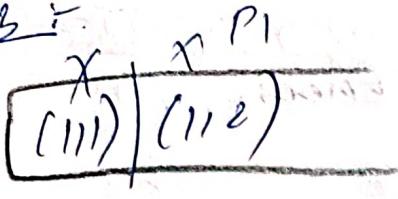
Recv Req (2,2)

Rep Rep (2,2)

The queue entry is deleted from its own when it enters C.S

Once reply is given, immediately Request is deleted, instead of waiting for Release msg.

No need of any explicit Release msg.
Reduced to 2 phase algo (Req, Reply)



Request (1,1)

8

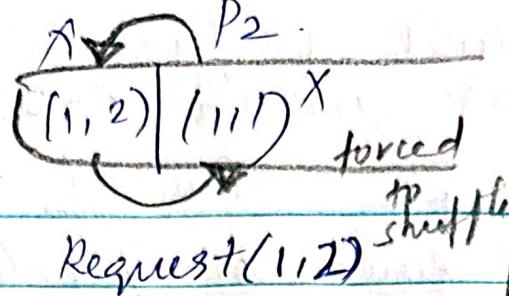
Recv Req (1,2)

Recv Rep(1,1): P₃

Recv Rep(1,1): P₂

exec C.S

Rep Req(1,2):



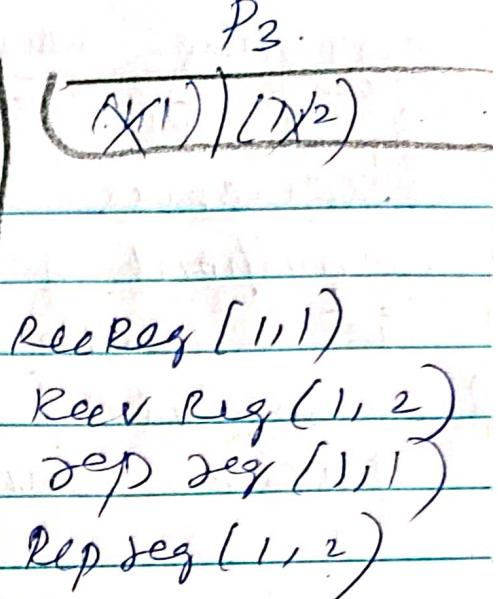
Request(1,2)

Recv Req (1,1)

Recv Rep(1,2): P₃

Recv Reply (1,2): P₁

Exec C.S



Recv Req (1,1)

Recv Rep (1,2)

Rep Seg (1,1)

Rep Seg (1,2)

new limit is a problem with fairness because an entry is deleted immediately a reply is sent.

↳ But this is solved bcoz,

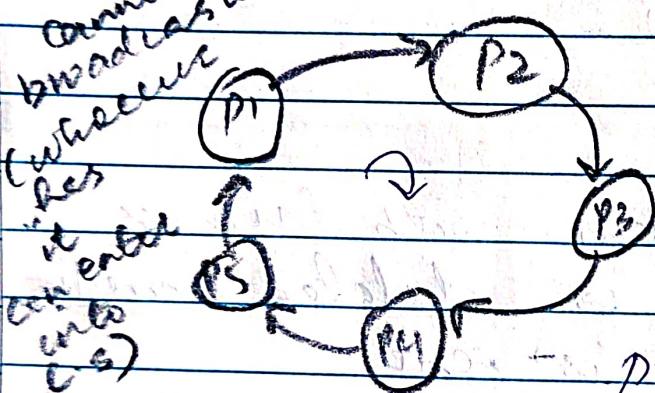
the entry is deleted in Recv, not in send. It is deleted only when it enters into C.S.

For outstanding requests, send new request after a time-out. By checking the sequence no. If it is older only it can process the higher live request, else it is an old Request.
 ↳ This solves liveliness problem.

Supporting mutual exclusion using contradiction
If 2 processes are in C-S, their process sequence no. should be same
(without happens condition), which is not possible.

so it is that only one process can exec in C-S

TOKEN BASED D-MUTEX (SUZUKI-KASAMI'S Algorithm) :-



Each process passes the token in the circular ring of processes which are predefined. Then if the process needs it, it takes the token, else passes it to neighbour.

Limitation:-

If P5 needs the token, P1 cannot send it to P2, not to P5.

The process does not know who needs the token. [does not know the demand]

This leads to starvation and has performance issue.

To solve this :-

2 Data structures are used :-

R - Request vector (available at every node)

T - Token Vector (available as only 1 total instance)
Size $\rightarrow n$ [no. of processes]

Algorithm :-

Request :-

Inc $R_i^o[i]$ [request for its own entry in R] and broadcast it

Response :- \rightarrow to identify outdated request, it should not process old request
on receiving the broadcast message,
 $R_j^o[i] = \max(R_j^o[i], R_i^o[i])$.
(every node)

If a node e has T then, if

$T[k] == R_j^o[i] - 1$, pass

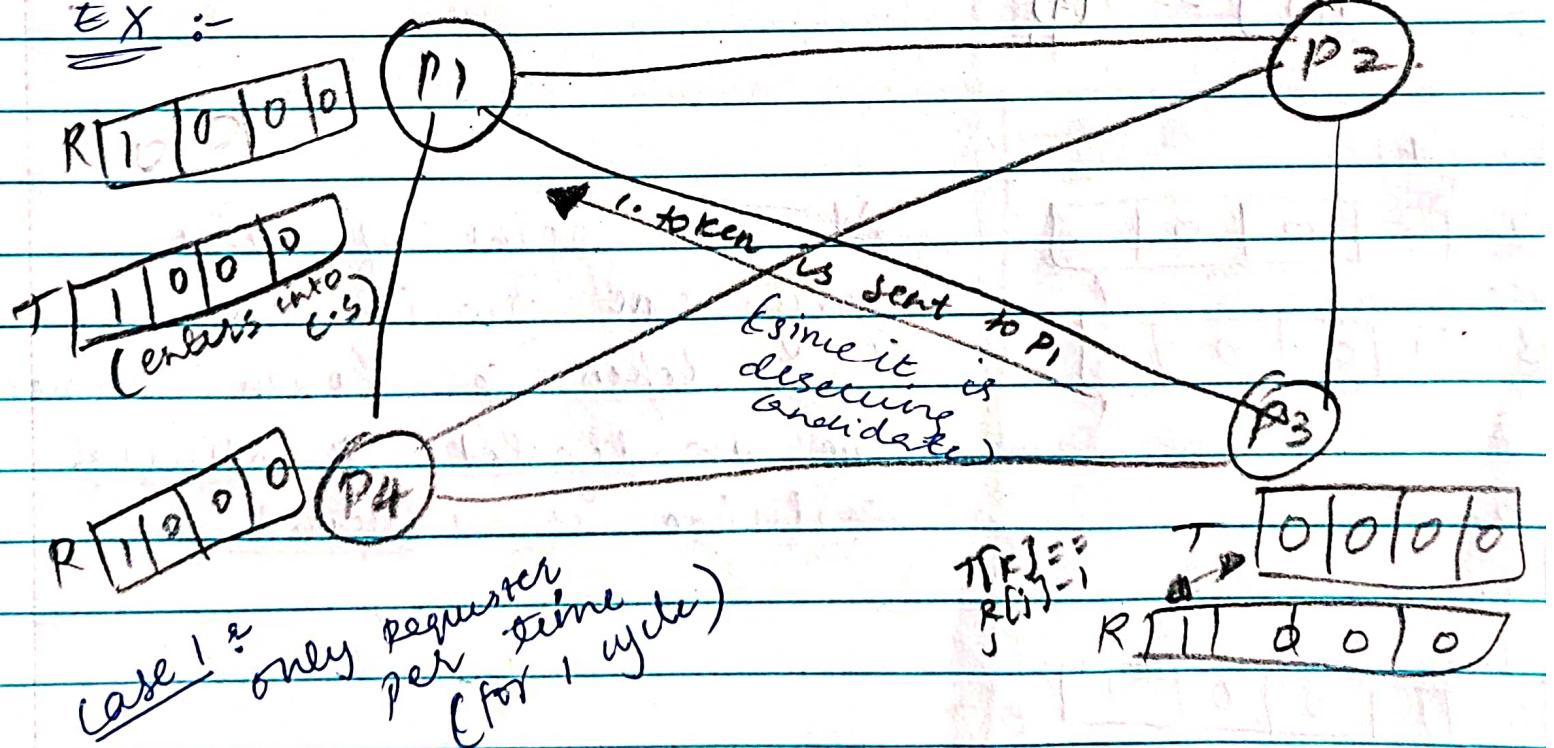
the token

$(k \rightarrow \text{for all values})$

entries]

$R[1|0|0|0]$

Ex :-



If token is not being used but it is possessed by a node \rightarrow IDLE TOKEN

In dictat ex:

P₁

1. R [0] 0 0 1 0

2. R [1] 0 1 0 1 0

3.

4.

5. T [0] 0 0 0 0

6. exec C.S

7. T [1] 0 1 0 1 0

8.

9. R [1] 0 0 1 1

10. T [1]

$$R[4] = T[4] + 1$$

P₄.

1. R [0] 0 1 0 1

2.

3. T [1] 0 0 1 0

4.

5.

8. T [1] 0 0 1 1

9.

10. T [Φ] 0 0 1 1

12. exec C.S

P₂

R [0] 0 1 0 1 0

R [1] 0 0 1 0

T [1]

$$R[3] = T[3] + 1$$

If a request has to be broadcasted, entry should be done earlier than the receiver's message. (start as seq nos. 1, 4) P between seq nos. 2, 3 id

[1] 0 1 0 1 1

[1] 1 0 0 1 1

13. T [1] 1 0 0 1 1

P₃

R [0] 0 1 0 1 0 0

T [0] 0 1 0 1 0

R [1] 1 0 1 0 1 0

T [1]

$$R[3] = T[3] + 1$$

If a request has to be broadcasted, entry should be done earlier than the receiver's message. (start as seq nos. 1, 4) P between seq nos. 2, 3 id

To achieve fairness :-

(Each process should remember its latest Request. When request is already pending, multiple requests cannot be made)

To achieve liveliness :-

only one in a order process is entering into the C.S, (Ex : First P₁ entry, next P₂, etc...)

Ex :-

P ₁	P ₂	P ₃
1. R ₁ [0 1 0 0]	R ₂ [0 1 0 0]	R ₃ [0 1 0 0] T T [0 1 0 0]
2. R ₁ [1 1 0 0] $R_i^o [i] = seq + 1$ broadcast (seq, P _i ^o)	R ₂ [1 1 0 0] . R ₂ [1 1 1 0] .	R ₃ [1 1 0 0] . T _j [k] = R _j [i] - 1 T _j [1 1 1 0]
3. T [1 0 0 0] exec C.S T [1 1 0 0]	T [1 1 0 0] exec C.S T [1 1 1 0]	R ₃ [1 1 1 0]
R ₁ [2 1 1 0]		R ₃ [1 1 1 1] ∞

concurrent requests

R₁ | 2 | 1 | 1 |

R₂ | α | 1 | 1 | .

R₃ | α | 1 | 1 |